

Dynamic Programming and MAB

ECE/CSE RL Monsoon 2021

September 16, 2021

✓ **Question 1.** Consider ten bandit arms numbered $1, 2, \dots, 10$. You pick an even numbered arm with probability twice that of picking an odd arm. Arm i gives a reward that is drawn from a Gaussian distribution with mean i and variance 1. You pick an arm with your above chosen policy 10 times. Derive from first principles the expected sum reward you receive at the end of picking ten times.

✓ **Question 2.** We have 10 arms. Arms 1, 2, 4, 5, 7, 9, 10 give a reward of 0 with probability 0.5 and a reward of 1 otherwise. The three other arms 3, 6, 8 give a reward of 0 with probability 0.3, a reward of 0.2 with probability 0.3, and a reward of 1 with probability 0.4. As always you want to maximize the expected reward. Derive six optimal stochastic policies.

✓ **Question 3.** Consider a variant of the ϵ -greedy policy where in we explore only over the non-greedy actions. Suppose that we have 3 bandit arms. Each arm gives a random reward that takes a value from the set $\{0, 1\}$. Assume an initial estimate $Q_1(a)$, $a \in \{1, 2, 3\}$. Choose the initial estimates to be different and non-zero for the arms. Create and explain an example sequence of $Q_t(a)$, A_t , R_t , for $t = 1, 2, 3, 4, 5, 6$, obtained under the assumption that you explore at odd times and exploit at even times. Use sample mean estimates.

✓ **Question 4.** Solve Exercise 3.4. Explain how you obtained the table. Your solution may be hand-written.

✓ **Question 5.** Write code that solves the linear equations required to find $v_\pi(s)$ and generate the values in the table in Figure 3.2. Note that the policy π picks all valid actions in a state with equal probability. Add comments to your code that explain all your steps.

✓ **Question 6.** Solve Exercises 3.15 and 3.16.

✓ **Question 7.** Write code that generates the optimal state-value function and the optimal policy for the Gridworld in Figure 3.5. You want to solve the corresponding system of non-linear equations. Explain all your steps.

✓ **Question 8.** Given an equation for v_* in terms of q_* .

✓ **Question 9.** Code policy iteration and value iteration (VI) to solve the Gridworld in Example 4.1. Your code must log output of each iteration. Pick up a few sample iterations to show policy evaluation and improvement at work. Similarly, show using a few obtained iterations that every iteration of VI improves the value function. Your code must include the fix to the bug mentioned in Exercise 4.4.

✓ **Question 10.** Code exercise 4.7.

✓ **Question 11.** When we defined a Markov Decision Process (MDP) we explicitly captured, using probability mass functions (PMF), the fact that the random variable R_{t+1} is dependent on the state S_t and the action A_t . Is the random variable R_{t+2} dependent on S_t and A_t ? Support your answer using the PMFs we used to define a MDP. [Hint: Start with the conditional PMF of R_{t+2} conditioned on S_t and A_t .]

✓ **Question 12.** Derive the expression for $E[R_{t+2}|S_t = s, A_t = a]$ in terms of the PMF(s) that define a MDP. [Hint: This is in a way an extension of Question 11.]

Question 13. We know that the state-value function $v_\pi(s) = E[G_t | S_t = s]$. Use this definition of $v_\pi(s)$ to derive the Bellman equation for $v_\pi(s)$ for all $s \in S$. The Bellman equation will use the PMF corresponding to the policy π and the PMF $p(s', r | s, a)$ and will provide a recursive method of calculating $v_\pi(s)$.

Question 14. Suppose an agent receives the sequence of rewards $R_1 = 2$, $R_2 = -1$, $R_3 = 10$, and $R_4 = -3$. Calculate the γ discounted return/reward for each time step for $\gamma = 0.5$. Also show that, if an agent receives a constant reward c , at every time step, for $\gamma < 1$, the infinite horizon discounted return is given by

$$G_t = \frac{c}{1 - \gamma}.$$

Question 15. Suppose you are given the optimal state-value function $v_*(s)$, for all $s \in S$. How will you find the optimal policy? Show your steps.

Question 16. A server requires information from a sensor. The server would like the information to be fresh. However, there is a cost to querying information from the sensor. Specifically, the state at the server can be either fresh or stale. The former indicates that the information at the server about the sensor is fresh and the latter indicates that it is stale. At any time, the server may choose to query or remain silent. A query makes stale information fresh with probability 0.8 and fresh information stale with probability 0.1. Staying silent keeps stale information stale with probability 1 and makes fresh information stale with probability 0.5. Also, a query when current information at the server is stale costs 8 dollars. Otherwise, it costs 4 dollars. Staying quiet has a reward of 4 dollars.

1. Draw and/or tabulate the MDP.

2. Consider a finite horizon problem in which the server gets to optimize costs over three time steps. Assume that if at the end of the three time steps the server has stale information it pays a cost of 10, else it receives a reward of 10. Also, assume that future costs are discounted with a factor $1/2$. Calculate the optimal costs for every starting state. Calculate the optimal policy. Show all your steps. Guesses and intuition will bring you unbounded costs.

3. Suppose the server is to optimize for ever and is looking for a stationary policy. Write down the first four iterations of value iteration that will help the server get closer to a stationary policy. Do the same for policy iteration (where one iteration includes evaluation and improvement). Clearly identify each iteration.

Question 17. Assume an infinite horizon discounted costs problem. Prove that the policy improvement step either improves the current policy or the current policy is optimal. Show all your steps and support them using the properties of dynamic programming.

Question 18. *The Stairwell:* A stairwell connects the ground floor (G) of a palace with its first floor (F). It consists of a total of n steps, indexed $1, 2, \dots, n$, with step 1 close to G and step n connected to F. A robot may start in any step and must learn to decide on whether it should go to F or to G. Both F and G are terminal states. The robot gets a reward of 2 with probability 0.5, and 0 otherwise, when it descends from step i to $i - 1$, $1 < i \leq n$. It also gets a reward of 1 on going from step 1 to G. On the other hand, the robot pays a cost of 1 (reward -1) when it moves from step i to $i + 1$, $1 \leq i < n$. Finally, the robot gets a bumper reward of $2n$ when it goes from step n to F. The robot can only take one step at a time.

Start with a policy in which the robot goes up or down a step in the stairwell with probability 0.5. Draw the corresponding MDP for $n = 2$. Use policy iteration to find the optimal policy for $n = 2$. Show all your iterations and explain them. How does the number of iterations increase with n ? Explain your answer. Assume a discount factor of 1.

Model for the Environment For All Following Questions

A person can be either *healthy* or *sick*. When healthy, a person may choose to either *stay home* or *go out*. If a healthy person chooses to stay home, the person stays healthy with probability 0.95 and falls sick with probability 0.05. A healthy person who stays healthy on staying home gets a random reward of 10 with probability 0.95 and 0 otherwise. A healthy person who falls sick on staying home gets a random reward of -10 with probability 0.9 and 0 otherwise.

A healthy person who chooses to go out, stays healthy with probability 0.7 and gets a random reward of 20 with probability 0.9 and 0 otherwise. A healthy person who chooses to go out falls sick with probability 0.3 and gets a random reward of -10 with probability 0.9 and 0 otherwise.

A sick person may either take medicine or choose not to take any medicine. A sick person who takes medicine stays sick with probability 0.1 and gets a deterministic reward of -2 . A sick person who takes medicine becomes healthy with probability 0.9 and gets a reward of -1 . A sick person who doesn't take medicine stays sick with probability 0.4 and gets a reward of -1 . A sick person who doesn't take medicine gets healthy with probability 0.6 and gets a reward of 0.

We will assume a discount factor of $\gamma = 0.9$ in all questions that follow.

Question 19. Draw the MDP.

Question 20. Consider a policy that in any state picks all allowed actions with equal probability. Derive the value of the policy.

In case you choose to follow an iterative method, write the values obtained at the end of each iteration. Otherwise, provide the final form of the system of equations you chose to solve and the final solution you obtained.

Question 21. Suppose at day t , a person is healthy and on day $t + 1$ the person falls sick. Derive the conditional expected return $E[G_t | S_t = \text{healthy}, S_{t+1} = \text{sick}]$. Assume the policy in the above question and feel free to use any results you derived in the previous question.

Explain in words (in a few lines) how you calculated the conditional expectation of R_{t+1} .

Question 22. Derive the optimal policy and the optimal value function using policy iteration. Calculate at most three iterations. If you find the optimal policy in fewer than three iterations, then just calculate as many iterations as required to find the optimal.

Provide the evaluation for each state and the improved policy obtained at the end of each iteration of policy iteration. So, in case you do three iterations, you must record $\pi_0, \pi_1, \pi_2, \pi_3$ and $v_{\pi_0}, v_{\pi_1}, v_{\pi_2}$ for all states.